Supervised Statistical Learning



Fabrizio Cipollini

DiSIA G. Parenti, Università di Firenze

MaBiDa - Florence, May-July, 2017



References

James, G., Witten, D., Hastie, T., and Tibshirani, R. (2014). *An Introduction to Statistical Learning: With Applications in R*. Springer Publishing Company, Incorporated.

- ▶ Background: Chapters 3 and 4 (and R)
- ▶ This course: Chapters 2, 5 (partially), 6, 7, 8
- Additional material provided during the course



Framework

Data: one dependent variable, p independent variables

$$(y_i,\underbrace{x_{i,1},\ldots,x_{i,p}}_{\mathbf{x}'_i})$$
 $i=1,\ldots,n$

- ▶ Models:
 - 1. **Quantitative** $y \rightarrow \text{Regression}$

$$egin{aligned} m{y} &= \mu(m{x}) + arepsilon \quad \text{s.t.} \quad arepsilon \perp m{x}, \ E(arepsilon) &= 0 \ y | m{x} \sim \left[\mu(m{x}), \sigma^2
ight] & \mu(m{x}) = E(y|m{x}) = f(m{x}) \end{aligned}$$

2. Categorical (0/1) $y \rightarrow$ Classification

$$y|oldsymbol{x} \sim Be(\mu(oldsymbol{x})) \qquad \mu(oldsymbol{x}) = E(y|oldsymbol{x}) = rac{e^{f(oldsymbol{x})}}{1 + e^{f(oldsymbol{x})}}$$

f(.) is unknown!



Challenge

 $\mathit{f}(\cdot)$ unknown $ightarrow \widehat{\mathit{f}}(\cdot)$ estimated using data

A challenge involving selection of

- Variables: f(·) may depend on a subset of the independent variables in the data (in principle, it could also depend on omitted variables!)
- ▶ Shape: For a given set of independent variables, each x_j may give a **linear** or a **non-linear** (...) contribution to $f(\cdot)$
- ► "Complexity": For a given set of independent variables, we can have **interactions** among the x_i 's into $f(\cdot)$
- All such aspects are interconnected!



Statistics vs Statistical Learning

	Statistics	Statistical Learning
Aspect	view	view
Slang	Conditional analysis	Supervised learning
Slarig	In/out-of sample	Training/Test sets
Objective	Inference	Prediction
Focus	Model interpretation	Prediction accuracy
	(parameters)	(error measures)
\widehat{f}	We need to know it	May be treated as black-box
Flexibility	Low	High

- Tools and methods are similar
- ► This course: Statistical Learning view



Definition and Decomposition

- ▶ Consider quantitative y: $y = f(x) + \varepsilon$, $\varepsilon \perp x$ and $E(\varepsilon) = 0$.
- Mean Squared Error:

$$\begin{split} \textit{MSE} := \textit{E}[(\textit{y} - \widehat{\textit{f}})^2] &= \underbrace{\textit{V}(\varepsilon)}_{|\textit{Irreducible}} &+ \underbrace{\textit{E}[(\textit{f} - \widehat{\textit{f}})^2]}_{|\textit{Reducible}} \\ &= \underbrace{\sigma^2}_{-/+\textit{E}(\widehat{\textit{f}})} &+ \textit{bias}^2(\widehat{\textit{f}}) + \textit{V}(\widehat{\textit{f}}) \end{split}$$

- $V(\varepsilon) = \sigma^2 =$ "Irreducible error"
- ▶ $bias(\hat{f}) = E(\hat{f}) f$ = "Error induced in approximating a (possibly) complex f by a (relatively) simple model"
- ► $V(\hat{f})$ = "How much \hat{f} would change if estimated using different data"
- ► Training MSE (MSE_{tr}) vs Test MSE (MSE_{ts})
- ► How to gain a better feeling? MBD2016-MSE-20160503.R



What we learn (theoretical point of view)

- Interpretation of the MSE and its components ("Irreducible Error", Variance, Bias²)
- 2. $\widehat{f}(\cdot)$ correctly specified (it includes all "terms" in $f(\cdot)$) \Longrightarrow Bias = 0
- Model complexity ↑ ⇒ Bias² ↓ and Variance ↑
 ⇒ Bias/Variance trade-off. An easy compromize?
- 4. No:
 - In simulations we know $f(\cdot)$ and we have many replications, so it is possible to compute Bias and Variance
 - ▶ In practice we cannot! (We have only one Training dataset)



What we learn (practical point of view)

5. We are interested in a small MSE_{ts} but we manage **one Training dataset**, guided by an estimate of the MSE_{tr} :

$$\widehat{MSE}_{tr} = \frac{1}{n} \underbrace{\sum_{i=1}^{n} [y_i - \widehat{f}(\mathbf{x}_i)]^2}_{SSR}$$

 \widehat{MSE}_{tr} (SSR) can be reduced over and over, by increasing model complexity... but if we exaggerate, this increases MSE_{ts} !

Thus, we need a compromise:

Fit **training data** as better as possible (low *SSR*), but . . .

don't go too far with model complexity!



Assembling Ingredients

► Recall our challenge:

Estimate the unknown f(x) using data

- To avoid indigestion, we split this lunch in three dishes:
 - 1. Variable Selection in a linear world → Chapter 6
 - 2. Shape Selection an additive world → Chapter 7
 - 3. "Complexity" Selection \rightarrow Chapter 8
- Always the target is the smallest MSE_{ts}, but we pursue this goal following a compromise strategy:
 Minimize SSR penalizing excessive model complexity

$$min PSSR = min(SSR + penalty(complexity))$$

In practice,

Penalized LS!



Introduction

• We assume a **linear** f(x):

$$f(\mathbf{x}) = \beta_0 + \beta_1 \mathbf{x}_1 \dots + \beta_p \mathbf{x}_p$$

(although we use the same symbol, here p can be lower that the number of variables in the data)

- Reasons for assuming linearity:
 - ► Can be a **good approximation** of the true f(x) (related to first order Taylor approximation of a math function)
 - We know what linear means while non-linear does not have a precise identity
 - By far more simple then more general approaches (think to OLS and the interpretation of coefficients)
 - Surprisingly competitive against more general approaches in many applications



Different flavors

- Subset selection:
 - Idea: Fit a lot of models based of different subsets of variables; then use some criterion to choose the best one
 - Variants:
 - Best subset
 - Stepwise: forward, backward, hybrid
- Shrinkage (regularization) methods:
 - Idea: Use all p independent variables, but penalize complexity by shrinking their coefficients toward zero
 - → Reduce Variance paying some price in terms of Bias²
 - Variants: (differing in the kind of penalty)
 - Ridge regression
 - Lasso regression
 - Elastic Nets
 - ▶ They work also with p > n!



Subset Selection: Best subset

- Algorithm:
 - 1. Start from the *null model* (no independent variables) \mathcal{M}_0
 - 2. For k = 1, ..., p
 - 2.1 Fit all models with exactly k predictors
 - 2.2 Select the best one and call it \mathcal{M}_k
 - 3. Select the best model among \mathcal{M}_k : k = 1, ..., p
- ▶ It requires p < n</p>
- ▶ # models = 2^p ($p = 20 \rightarrow \#$ models > 1000000)



Subset Selection: Forward Stepwise

- Algorithm:
 - 1. Start from the *null model* (no independent variables) \mathcal{M}_0
 - 2. For k = 0, ..., p 1
 - 2.1 Fit all p k models increasing \mathcal{M}_k by 1 predictor
 - 2.2 Select the best one and call it \mathcal{M}_{k+1}
 - 3. Select the best model among \mathcal{M}_k : $k = 1, \dots, p$
- ▶ It can be used also when p > n
- ▶ # models = p(p+1)/2 + 1 ($p = 20 \rightarrow \#$ models = 211)



Subset Selection: Backward Stepwise

- Algorithm:
 - 1. Start from the *full model* (all independent variables) \mathcal{M}_p
 - 2. For k = p, ..., 1
 - 2.1 Fit all k models decreasing \mathcal{M}_k by 1 predictor
 - 2.2 Select the best one and call it \mathcal{M}_{k-1}
 - 3. Select the best model among \mathcal{M}_k : $k = 1, \dots, p$
- ▶ It requires p < n</p>
- ▶ # models = p(p+1)/2 + 1 ($p = 20 \rightarrow \#$ models = 211)



Subset Selection: Hybrid Stepwise

- Algorithm: intermediate between forward and backward
- Idea:
 - In a forward based algorithm: after adding variables it may go back removing variable not improving fit
 - In a backward based algorithm: after removing variables it may go forward adding variables improving fit



Subset Selection: Selecting the best

Minimize Akaike or Bayesian Information Criteria

$$AIC = SSR + 2k\widehat{\sigma}_*^2$$
$$BIC = SSR + \ln(n)k\widehat{\sigma}_*^2$$

- ▶ k = #parameters
- $\hat{\sigma}_*^2 = \text{estimate of } \sigma^2 \text{ based on a "large" (?) model}$
- Comments:
 - Structure: SSR + penalty(k)
 - Formulas proportional to those in the book (pp. 211-212)
 - Further criteria in the book (not in step())
- Commonly used versions are different (cf step()):

$$AIC = n \ln(SSR) + 2k + const$$

 $BIC = n \ln(SSR) + \ln(n)k + const$ $const = n (\ln(2\pi/n) + 1)$



Subset Selection: Let's Eat!

- Appliances Energy Consumption data: energy use of appliances in a low energy building
- Dependent variable: energy use of appliances
- Independent variables:
 - temperature and humidity measures from sensors of different building spaces
 - weather from a nearby airport station
 - recorded energy use of lighting fixtures
- ▶ n = 19737 vs p = 33 predictors
- ▶ MBD2017-Energy-20170513.R



Shrinkage Methods: Ridge Regression

Estimate parameters minimizing

$$SSR + \lambda \frac{1}{2} \sum_{j=1}^{p} \beta_j^2$$
penalty

- lacktriangleright $\lambda \geq 0$ is the penalty/shrinkage/regularization parameter
- lacksquare $\lambda=0$ (o OLS) vs $\lambda=\infty$ (o $\widehat{eta}_1=\ldots=\widehat{eta}_p=0$)
- ► Further comments on: method idea, *L*₂-norm, 1/2 factor
- Not scale equivariant: standardize predictors!
- Large gain when OLS have large variance: large p; highly correlated predictors
- Very fast: computations required to estimate parameters for all λ's ≈ OLS computations
- ► Let's try: MBD2017-Energy-20170513.R



Ridge Regression: How to Choose λ ?

- Use Cross Validation (CV, chapter 5), an approach that can be applied to almost any statistical learning method
- ► *k*-fold CV:
 - ▶ Randomly divide data in k groups (typically k = 5, 10, 20) of approximately equal size ($\approx n/k$)
 - ▶ For j = 1, ..., k
 - Use the j-th group as test set and the remaining as training set
 - Estimate using the training set
 - ▶ Use it to predict y in the j-th and compute MSE_j
 - Then

$$CV_{(k)} = \frac{1}{k} \sum_{j=1}^{k} \widehat{MSE}_{j}$$

- ▶ **Leave One Out CV** (LOOCV): Special case obtained for k = n (each group has one observation) \rightarrow comments
- ► Let's finish: MBD2017-Energy-20170513.R
 Cioollini Supervised Statistical Learning



Shrinkage Methods: Lasso Regression

 Only a "small" difference with Ridge Regression: Estimate parameters minimizing

$$SSR + \lambda \sum_{j=1}^{p} |\beta_j|$$
penalty

- lacksquare $\lambda=0$ (o OLS) vs $\lambda=\infty$ (o $\widehat{eta}_1=\ldots=\widehat{eta}_p=0$)
- L₁ penalty instead of L₂ penalty
- ▶ Why L_1 in place of L_2 ? → Sparsity



Shrinkage Methods: Lasso vs Ridge (1)

Lasso

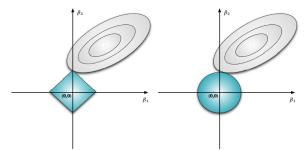
$\min\left(\mathit{SSR} + \lambda \sum_{j=1}^p |\beta_j|\right)$

$$\min(SSR)$$
 s.t. $\sum_{i=1}^{n} |\beta_j| \leq s$

Ridge

$$\min\left(SSR + \lambda \frac{1}{2} \sum_{j=1}^{p} \beta_j^2\right)$$

$$\min(SSR) \text{ s.t. } \sum_{j=1}^{j=1}^{p} |\beta_j| \le s \qquad \min(SSR) \text{ s.t. } \frac{1}{2} \sum_{j=1}^{p} \beta_j^2 \le s$$





Shrinkage Methods: Lasso vs Ridge (2)

	Lasso	Ridge
Sparsity (parsimony)	Yes	No
Interpretation	High	Low
Smooth coefficients path	No	Yes
Selection with a group of	At most one predictor	All
highly correlated predictors	(does not care which one)	All
Performance with $p < n$ and	Ridge tends to win	
highly correlated predictors	Hidge terias to will	
Selection with $p > n$	At most <i>n</i> predictors	All

and so?



Shrinkage Methods: Elastic Nets (1)

Estimate parameters minimizing

$$SSR + \lambda \left((1 - \alpha) \frac{1}{2} \sum_{j=1}^{p} \beta_j^2 + \alpha \sum_{j=1}^{p} |\beta_j| \right)$$
penalty

where $\alpha \in [0, 1]$

Equivalent to

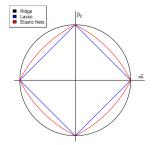
min(*SSR*) s.t.
$$(1 - \alpha)\frac{1}{2}\sum_{j=1}^{p}\beta_{j}^{2} + \alpha\sum_{j=1}^{p}|\beta_{j}| \leq s$$

(see next slide)



Shrinkage Methods: Elastic Nets (2)

Parameter's feasible set:

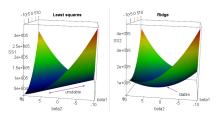


- Motivations:
 - ► The L₁ part of the penalty generates sparsity
 - ► The L₂ part of the penalty
 - ▶ Removes the limitation on the **number** of selected variables
 - Encourages grouping effect
 - Stabilizes the coefficients path



Names

Ridge



In case of multicollinearity, the criterion function (SSR) has a long valley, corresponding to a **ridge** if -SSR is considered.

Ridge regression "fixes" the ridge, adding a penalty that turns the ridge into a nice peak

- LASSO = Least Absolute Shrinkage and Selection Operator
- Elastic Net = Elastic net between the Ridge and the Lasso



Introduction

We move from a linear "world"

$$f(\mathbf{x}) = \beta_0 + \beta_1 \mathbf{x}_1 + \ldots + \beta_p \mathbf{x}_p$$

to a possibly non-linear one

$$f(\mathbf{x}) = \beta_0 + f_1(x_1) + \ldots + f_p(x_p)$$

while retaining additivity of each variable

We start from the case p = 1

$$y = f(x) + \varepsilon$$

How to structure the (possibly) non-linear $f(\cdot)$? Since $f(\cdot)$ is unknown, we need a *flexible* but *simple* way to **represent** it.



How to Represent $f(\cdot)$?

1-st idea: polynomial

$$f(\mathbf{x}) = \beta_0 + \beta_1 \mathbf{x} + \ldots + \beta_k \mathbf{x}^d$$

Pros: simple, easy interpretation *Cons*:

- Numerically unstable for high exponents
- Global structure: polynomials fit well in regions where there are lot of data but take "crazy" shapes elsewhere
- ▶ 2-nd idea: piecewise polynomials → splines
 - Divide the x range in intervals, separated at boundary points (→ knots)
 - ► Use a low degree polynomial in each interval (→ typically, d = 1 or d = 3)
 - Impose some kind of smoothness at boundary points

Pros: numerical stability, local structure



Cubic Splines

- **Definition**: 3-rd degree piecewise polynomials C^2 at knots
- ▶ Consider K knots $\rightarrow \xi_1 < \ldots < \xi_K$
- How many free coefficients?

$$\underbrace{4(K+1)}_{\text{number of parameters}} - \underbrace{(K+K+K)}_{C^2 \text{ continuity}} = K+4$$

Direct math representation incorporating knot constraints:

$$f(x) = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 (x - \xi_1)_+^3 + \ldots + \beta_{K+4} (x - \xi_K)_+^3$$

Important: linear in the parameters!

 There are computationally more efficient representations (although more cumbersome to write). Linearity in the parameters is preserved



Natural Cubic Splines

- Outside the external knots there are usually few data \rightarrow Variance of $\hat{f}(\cdot)$ large here \rightarrow How to remedy?
- Impose boundary constraints: linearity outside the external knots (equivalent to 2-nd and 3-rd derivatives equal to zero)
- ▶ Natural Cubic Splines = cubic splines + boundary constraints
- Number of free coefficients

$$\underbrace{4(K+1)}_{\text{number of parameters}} - \underbrace{(K+K+K)}_{C^2 \text{ continuity}} - \underbrace{(2+2)}_{\text{boundary contraints}}) = K$$

- Direct math representation incorporating knot constraints not simple to write. Anyway, **linearity** in the parameters is retained
- Hereafter natural cubic splines! Cipollini



How Many Knots? And Where?

Two approaches:

- ► Regression Splines: *K* << *n* and knots placed at "strategic" positions
- ▶ Smoothing Splines: knots at unique x values (ties removed), so that K = n (or $\approx n$ with ties)



Regression Splines

For given K (<< n) and **location** of knots, it is similar to linear regression (coefficients estimated by minimizing SSR)

- ► How much *K*?
 - ▶ K large \Rightarrow more flexible $\widehat{f}(\cdot) \rightarrow small$ Bias², large Variance K small \Rightarrow less flexible $\widehat{f}(\cdot) \rightarrow large$ Bias², small Variance
 - How to find a good compromise? Compare different K's using k-fold CV
- Where to place knots (for a given K)?
 - Equally spaced values over the x range
 - Equally spaced quantiles of x (without ties)
 - Adaptive schemes (usually not employed)
- ▶ Then, minimize

$$SSR = \sum_{i=1}^{n} [y_i - s(x_i)]^2$$
 $s(x) = \beta_1 s_1(x) + \ldots + \beta_K s_K(x)$



Smoothing Splines

Knots are placed at all unique x values but the SSR is penalized by a term related to the curvature of \hat{f}

Then, minimize

$$PSSR = \sum_{i=1}^{n} [y_i - s(x_i)]^2 + \lambda \int (s''(t))^2 dt$$
$$s(x) = \beta_1 s_1(x) + \ldots + \beta_n s_n(x)$$

Remark: natural cubic splines minimize

$$\sum_{i=1}^{n} [y_i - g(x_i)]^2 + \lambda \int (g''(t))^2 dt$$

for any $g(\cdot)!$

Remark: n coefficients ... but they are not free!



λ and Degrees of Freedom

- Smoothing splines have n coefficients, but they are not free. We want to derive an "equivalent number of free coefficients" = degrees of freedom = df
- In linear models, fitted values can be computed as

$$\widehat{\boldsymbol{y}} = \underset{(n,1)}{\boldsymbol{X}} \widehat{\boldsymbol{\beta}} = \underbrace{\boldsymbol{X}(\boldsymbol{X}'\boldsymbol{X})^{-1}\boldsymbol{X}'}_{(n,n)} \underbrace{\boldsymbol{y}}_{(n,1)} = \boldsymbol{H}\boldsymbol{y}$$

df = p = # of columns of X = trace(H)

Same approach in smoothing splines:

$$\widehat{m{y}} = m{H}_{\lambda} \ m{y}$$
 $(n,1) = (n,n)(n,1)$
 $df = \operatorname{trace}(m{H}_{\lambda})$

$$\lambda \uparrow \Leftrightarrow df \downarrow$$

But how to choose λ or df?



Choosing λ or df

Use CV, in particular LOOCV based on SSR

$$LOOCV_{\lambda} = \sum_{i=1}^{n} [y_i - \underbrace{\widehat{y}_i^{(-i)}}_{\text{depends on } \lambda}]^2 = \sum_{i=1}^{n} \left[\frac{y_i - \widehat{y}_i}{1 - \{H_{\lambda}\}_{i,i}} \right]^2$$

Also named **PRESS** (Predicted Residual Error Sum of Squares) statistic

- Computationally simple formula, based on the original fit to all data. This happens any time the model is linear in the parameters
- ▶ In practice: compute $LOOCV_{\lambda}$ for different choices of λ (or df) then select the best one



Additive Models

Go back to the complete model

$$y = f(\mathbf{x}) + \varepsilon$$
 $f(\mathbf{x}) = \beta_0 + f_1(x_1) \dots + f_p(x_p)$

- ▶ Use the approach described for any $f_i(\cdot)$ (j = 1, ..., p).
- ▶ To identify β_0 and the $f_j(\cdot)$'s we impose

$$\sum_{i=1}^{n} f_j(x_{i,j}) = 0 \quad j = 1, \dots, p \quad \Rightarrow \quad \widehat{\beta}_0 = \overline{y}$$

- $ightharpoonup \widehat{f}_j(\cdot)$'s are estimated by **backfitting**. For $j=1,\ldots,p$:
 - $\widehat{f}_{j}(\cdot)$ estimated by using partial residuals (residuals on all independent variables but x_{i}) as dependent variable
 - The procedure is iterated until convergence
- ► Snack time! MBD2017-Energy-GAM-20170513.R



How to test for linearity of one component

- Fit the model under analysis but with the variable under testing (say x) taken linear (in R: x instead of s (x))
- ► F-test (D = Residual Deviance = SSE; 0 = linear model)

$$F = \frac{\frac{D_0 - D}{df - df_0}}{\frac{D}{n - df}} | H_0 \approx F(df - df_0, n - df_0)$$

or

$$F = rac{rac{D_0 - D}{rdf_0 - rdf}}{rac{D}{rdf}} | H_0 pprox F(rdf_0 - rdf, rdf_0)$$

 $(rdf = \text{`residual df'} = n - df, \text{ see the } \mathbb{R} \text{ output)}$



Introduction

- Aim: To build a predictive model without too much assumptions regarding
 - which variables
 - their corresponding functional form
 - their interaction
- Basic idea: To resort to iterative/adaptive methods, trying to capture complexity where it is stronger
 - 1. **Start** from a "**null**" model (no independent variables)
 - 2. "Add" the contribution of the best fitting variable in the data
 - "Add" the contribution of the best fitting variable in the data, possibly in interaction with the variables already included
 - 4. Continue as in 3 until some stopping criterion is satisfied
 - (Possibly) Since the increase in "complexity" (2-4) can happen in a suboptimal way, the model in 4 is **pruned** to reach a reasonable compromise between fit and complexity



Recursive Partitioning

- Recursive Partitioning means Recursive Binary Splitting
- Procedure:
 - 1. Start from a "**null**" model (no predictors)
 - Find x_j and the cutpoint s which "best" split the data into two regions

$$R_1^{(L)} = \{i : x_{i,j} < s\}$$
 $R_1^{(R)} = \{i : x_{i,j} \ge s\}$

"Best" means to minimize SSR:

$$\sum_{i \in R_1^{(L)}} \left(y_i - \overline{\boldsymbol{y}}_{\boldsymbol{R}_1^{(L)}} \right)^2 + \sum_{i \in R_1^{(R)}} \left(y_i - \overline{\boldsymbol{y}}_{\boldsymbol{R}_1^{(R)}} \right)^2$$

- 3. Find x_j and the cutpoint s which "best" splits **one region** into two sub-regions
- 4. Continue as in 3 until some stopping criterion is satisfied
- The model built can be too complex; use CV to trim back the full tree



Recursive Partitioning: Final Product

M exhaustive non-overlapping regions ("rectangles")

$$R_m : m = 1, ..., M$$

Each i belongs to only one R_m

In each region the **prediction** is constant:

$$\overline{\mathbf{y}}_{R_m} = \operatorname{mean}\{y_i : i \in R_m\}$$

Such regions minimize

$$SSR = \sum_{m=1}^{M} SSR_m = \sum_{m=1}^{M} \sum_{i \in R_m} (y_i - \overline{y}_{R_m})^2$$

► The final product can be visualized by a binary tree where knots are characterized by (variable, cutpoint, left/right)



Recursive Partitioning: the Name

- Recursive Partitioning is sometimes known with as CART (Classification And Regression Trees)
- A curiosity:
 - CART is trademarked
 - tree has been used for an S-Plus package (also available in R)
 - The name Recursive PARTitioning (rpart) was chosen for the R package. Today, rpart is more common than the original CART...the power of free software!
- ▶ Let's taste: MBD2017-Energy-rpart-20170606.R



Introduction

- Topic not covered in the book
- So far: We aimed at predicting y using information in x
- ▶ Reference formulation: $y = f(x) + \varepsilon$, $\varepsilon \perp x$ and $\varepsilon \sim [0, \sigma^2]$
- Although interested in predicting y using predictors x, there may be reasons to model a transformation of y

$$w = f(x) + \varepsilon$$
 $\varepsilon \perp x$ and $\varepsilon \sim [0, \sigma^2]$

with w = g(y) 1-to-1. To go back, $y = g^{-1}(w)$

Examples:

<i>y</i>	$g(\cdot)$	$g^{-1}(\cdot)$
Energy	$w = \ln y$	$y = \exp(w)$
$\in (0, \infty)$	$w = \sqrt{y}$	$y = w^2$
N. medical examinations	$w = \ln y$	$y = \exp(w)$
\in [1, ∞)	$w = \sqrt{y-1}$	$y=w^2+1$



Motivations of and Issues

- Why transform?
 - Model diagnostics on w can look better than for y: for example, homoskedastic and/or less skewed residuals
 - Predictions of y can be better passing through w than directly modeling y
- Issues:
 - ► Goodness of fit statistics (R², SSR, AIC and similar) on different transformations of y cannot be compared
 - ▶ An unbiased **prediction** of *y* **is not** simply $g^{-1}(\widehat{w})$:

$$\widehat{y} = E(y|\mathbf{x}) \underset{y=g^{-1}(w)}{=} E(g^{-1}(w)|\mathbf{x})$$

$$\neq g^{-1}(\cdot) \text{ non linear}$$

$$g^{-1}(\widehat{w})$$

So? How to get \hat{y} ?



Predicting *y* when Modeling w = g(y)

Exact (case by case)

W	ŷ	Assumption
${\sqrt{y}}$	$\widehat{\mathbf{w}}^2 + \widehat{\sigma}^2$	$arepsilon \sim (0, \sigma^2)$
$\sqrt{y-1}$	$\widehat{w}^2 + \widehat{\sigma}^2 + 1$	$arepsilon \sim (0, \sigma^{2})$
ln(y)	$e^{\widehat{w}}\widehat{E}(e^{arepsilon})$	$arepsilon \sim (0, \sigma^2)$
ln(y)	$oldsymbol{e}^{\widehat{oldsymbol{w}}}oldsymbol{e}^{\widehat{\sigma}^2/2}$	$arepsilon \sim \textit{N}(0, \sigma^2)$
$\ln(y/(1-y))$?	?

▶ 2-nd order **Taylor expansion** of $g^{-1}(w)$ around \widehat{w}

$$\widehat{y} pprox g^{-1}(\widehat{w}) + g^{-1}(\widehat{w})\widehat{\sigma}^2/2$$

Average using residuals

$$\widehat{y} = \frac{1}{n} \sum_{i=1}^{n} g^{-1} (\widehat{w} + \widehat{\varepsilon}_i)$$



Introduction

Statistical Learning	Statistics	Example
Regression	y quantitative	expenditure
negression	(continuous)	$\in [0,+\infty)$
Classification	y qualitative	buy?
Classification	(categorical)	$\in \{\textit{no}, \textit{yes}\}$

- We could have more than two categories (example?)
- In case of more categories, they can be ordered or not (example)? Models for orderer and unordered categories differ
- Hereafter, we focus on models with two categories, identified as 1 and 0



Modeling Framework (Essential Ideas)

	Distribution	Link	(Linear)
			Predictor
Regression	$m{y} m{x} \sim m{N}(\mu(m{x}), \sigma^2)$	$\mu = \eta$	$\eta(\mathbf{x}) = \mathbf{x}'\boldsymbol{\beta}$
Classification	$m{y} m{x}\sim m{Be}(\mu(m{x}))$	$\mu = rac{ extbf{e}^{\eta}}{ extbf{1} + extbf{e}^{\eta}}$	$\eta({m x})={m x}'{m eta}$

- ► The approach can be extended to different distributional assumptions (for example to model counting data)
 → family of Generalized Linear Models (GLM)
- Link functions are monotone
- The linear predictor can be replaced by non-linear formulations in x (example GAM)
- ▶ **Penalizing** β , we can get Ridge, Lasso or Elastic Nets versions of GLM's
- Estimated by Maximum Likelihood (GLM) or their penalized counterparts (GAM, Ridge, Lasso, Elastic Nets)



Predictions

How to transform
$$\widehat{\mu}(\boldsymbol{x}) = \widehat{P}(y=1|\boldsymbol{x}) \in [0,1]$$
 to $\widehat{y} \in \{0,1\}$?

General classification rule:

$$\widehat{y} = 1 \Leftrightarrow \widehat{\mu}(x) > c$$

► How to set the cut-off c?



Setting the cut-off

- ▶ **Typical** classification rule sets c = 0.5. It makes sense only when 0/1 in the data are balanced.
- ▶ **Uninformed** classification rule sets $c = \widehat{\mu}$. It takes the *unconditional* probability of 1, $\widehat{\mu}$, as reference. It may lead to *biased* in sample classification, in the sense $\sum_{i=1}^{n} \widehat{y}_i \neq \sum_{i=1}^{n} y_i$
- ▶ **Unbiased** classification rule sets $c = \text{quantile}_{1-\widehat{\mu}}(\widehat{y})$. It implies $\sum_{i=1}^{n} \widehat{y}_{i} \simeq \sum_{i=1}^{n} y_{i}$
- Minimum Cost rule sets c to minimize the miss-classification cost → next slide



Minimum Cost Rule

		$\widehat{m{y}}$		
У	$\widehat{P}(y=* \mathbf{x})$	0	1	
0	$1-\widehat{\mu}(oldsymbol{x})$	<i>c</i> ₀₀	<i>c</i> ₀₁	
1	$\widehat{\mu}(oldsymbol{x})$	<i>c</i> ₁₀	<i>C</i> ₁₁	
	$\widehat{E}(cost \widehat{y}=*)$	$egin{aligned} c_{00}(1-\widehat{\mu}(oldsymbol{x}))\ +c_{10}\widehat{\mu}(oldsymbol{x}) \end{aligned}$	$c_{01}(1-\widehat{\mu}(oldsymbol{x})) \ +c_{11}\widehat{\mu}(oldsymbol{x})$	

$$\begin{split} \widehat{y} &= 1 \Leftrightarrow \widehat{E}(cost|\widehat{y} = 1) \leq \widehat{E}(cost|\widehat{y} = 0) \\ &\Leftrightarrow \widehat{\mu}(\mathbf{x}) \geq \frac{1}{1 + \frac{c_{10} - c_{11}}{c_{01} - c_{00}}} \underbrace{=}_{c_{00} = c_{11} = 0, c_{10} = 1} \frac{c_{01}}{1 + c_{01}} \end{split}$$

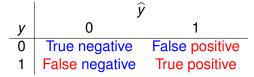
Particular cases:

- If $c_{10} = 1 \rightarrow \widehat{\mu}(\mathbf{x}) \geq 0.5$ (typical rule)
- ▶ If $c_{10} = \frac{\widehat{\mu}}{1-\widehat{\mu}} \to \widehat{\mu}(\mathbf{x}) \ge \widehat{\mu}$ (uninformed rule)
- ▶ If $c_{10} = \frac{qnt}{1-qnt} \rightarrow \widehat{\mu}(\mathbf{x}) \geq qnt$ (unbiased rule)



Evaluating Predictions: Confusion Matrix

Cells and names



 Using a classification rule, one can get a matrix of frequencies by cell

	ĺ		
У	0	1	
0	<i>n</i> ₀₀	<i>n</i> ₀₁	<i>n</i> _{0.}
1	<i>n</i> ₁₀	<i>n</i> ₁₁	$n_{1.}$
	<i>n</i> _{.0}	n _{.1}	n



Evaluating Predictions: Statistics

The matrix of frequencies

	ĺ		
У	0	1	
0	n_{00}	<i>n</i> ₀₁	$n_{0.}$
1	<i>n</i> ₁₀	n ₁₁	$n_{1.}$
	<i>n</i> .0	n _{.1}	n

- ► Classification cost: $n_{10} + n_{01}c_{01}$
- ► Classification accuracy: $(n_{00} + n_{11})/n$
- ► Classification error: $(n_{10} + n_{01})/n$
- ► (*) **Sensitivity** (True Positive Rate):

$$\hat{P}(\hat{y}=1|y=1)=n_{11}/n_{1.}$$

Specificity (True Negative Rate):

$$\widehat{P}(\widehat{y} = 0 | y = 0) = n_{00}/n_{0.}$$

► (*) **False Positive Rate** (1-Specificity):

$$\widehat{P}(\widehat{y} = 1 | y = 0) = n_{01}/n_{0.} = 1 - n_{00}/n_{0.}$$



Evaluating Predictions: The ROC Curve

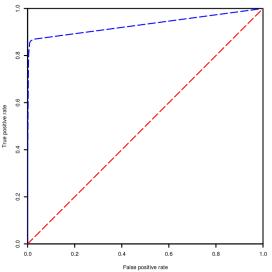
Plot of
$$(1 - Specificity, Sensitivity) = (FPR, TPR)$$

when the **cut-off** c **moves** from 0 to 1

- To understand: on a blank graph, report (FPR, TPR) for the following key configurations
 - $c = 0 \rightarrow ?$
 - ▶ $c = 1 \rightarrow ?$
 - ▶ Perfect predictions (no errors) for any $c \rightarrow ?$
 - ▶ Purely random predictions for any $c \rightarrow$?
- ► To paint: apply the definition ROCPION
- ► To **synthesize** (graphical ideas ROC Plot):
 - ► AUC ∈ [0.5, 1]
 - ▶ Accuracy Ratio = Gini = $2 \cdot AUC 1 \in [0, 1]$



Evaluating Predictions: ROC Plot



ROC Definition



Further Error/Performance Measures: R²

$$R^2 = \frac{dev(REG)}{dev(y)} \in [0, 1]$$

It is based on the deviance decomposition

$$dev(y) = dev(REG) + dev(RES)$$

But what such dev's? No more as in the linear model. . .

▶ I_0 (null log-lik) ≤ I (model log-lik) ≤ I_S (saturated log-lik)

$$dev(y) = 2(I_S - I_0)\sigma^2$$

$$dev(REG) = 2(I - I_0)\sigma^2 \qquad dev(RES) = 2(I_S - I)\sigma^2$$

► For each GLM member, such *dev*'s have specific formulas



Further Error/Performance Measures: Cross-Entropy

▶ With a 0/1 dependent,

$$dev(RES) = 2 \cdot \underbrace{\left(-\sum_{i=1}^{n} \left[y_{i} \ln \mu(\mathbf{x}_{i}) + (1 - y_{i}) \ln(1 - \mu(\mathbf{x}_{i}))\right]\right)}_{n \cdot \text{Cross-Entropy}}$$

Then

$$CE = -\frac{1}{n} \sum_{i=1}^{n} \left[y_i \ln \mu(\mathbf{x}_i) + (1 - y_i) \ln(1 - \mu(\mathbf{x}_i)) \right] \in [0, \ln 0.5]$$



Classification Trees vs Regression Trees

- Regression Trees are used y is quantitative
- ► Classification Trees are very similar, with the difference that the dependent variable is **qualitative** (0/1)
- ▶ **Criterion** to increase/prune the tree: denoting $\hat{p}_m/\hat{q}_m =$ estimated proportion of 1/0 in R_m ,
 - Increasing based on the sum of

$$G_m = \widehat{p}_m(1 - \widehat{p}_m) + \widehat{q}_m(1 - \widehat{q}_m) = 2\widehat{p}_m\widehat{q}_m$$
 Gini Index $D_m = -\left[\widehat{p}_m\log\widehat{p}_m + \widehat{q}_m\log\widehat{q}_m\right]$ Cross-Entropy

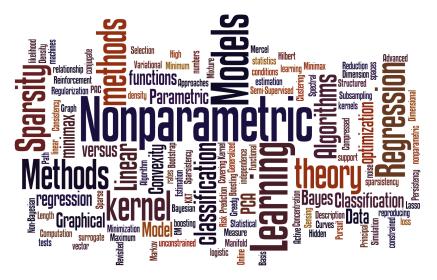
Pruning based on the sum of

$$E_m = 1 - max\{\hat{p}_m, \hat{q}_m\}$$
 Classification Error Rate

For everything else, they work very similarly



Today





Tomorrow

